



DAPO Math · Qwen3.5-9B-Base · fixed 8×H200 trainer + 8×H200 inference

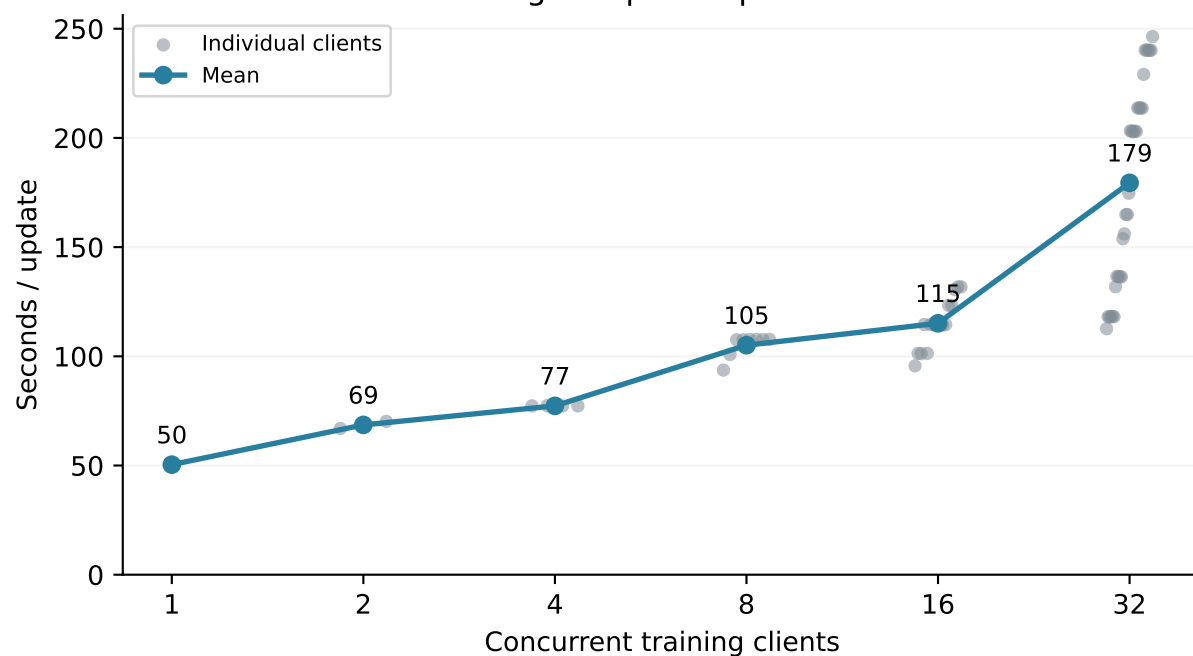
Total useful output throughput



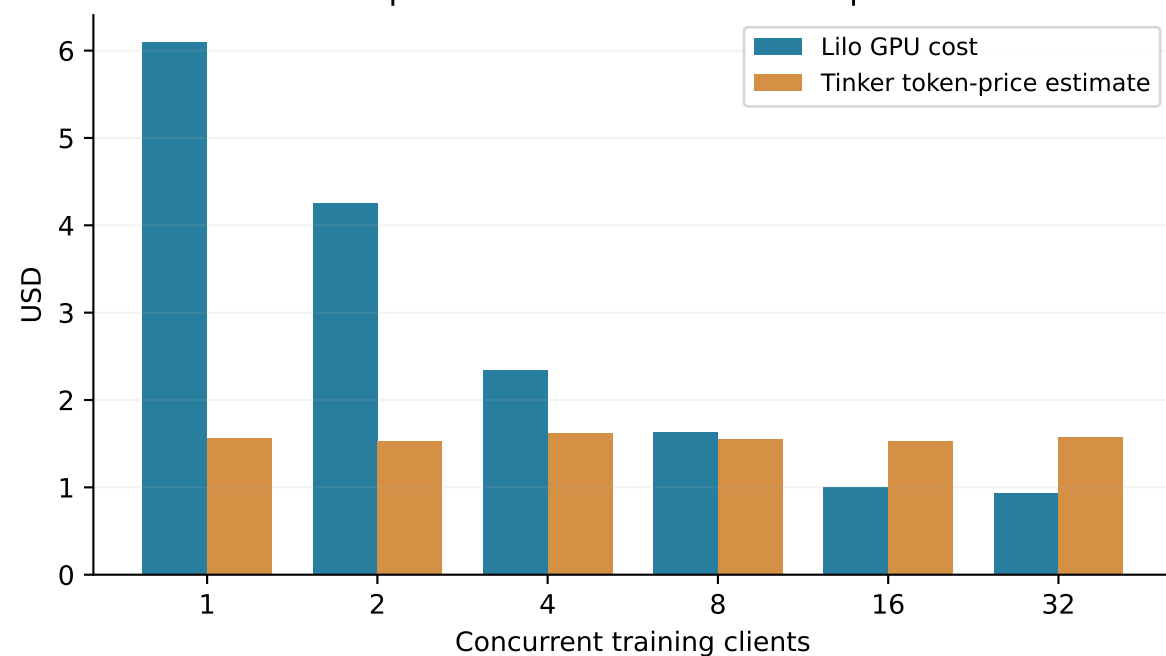
Throughput per client (total ÷ clients)



Average step time per client



Cost per client for six measured updates



8 updates/client; first 2 excluded. Async within measured window. No checkpoints/evaluations.
 Lilo: GPU-only wall cost. Tinker: same tokens, 7/8 prompt copies assumed cached; not a Tinker run.